

CLaC at SemEval-2026 Task 6: Response Clarity Detection in Political Discourse

Nawar Turk, Lucas Miquet-Westphal, Leila Kosseim

Computational Linguistics at Concordia (CLaC) Lab

Dept. of Computer Science and Software Engineering

Concordia University, Montréal, Québec, Canada

{nawar.turk, lucas.miquet-westphal}@mail.concordia.ca,
leila.kosseim@concordia.ca

Abstract

In this paper, we present our system for SemEval-2026 Task 6 (CLARITY) on response clarity and evasion detection in question-answer pairs from U.S. presidential interviews, comparing fine-tuned encoders with prompt-based LLMs. Our LLM ensemble achieves 80 macro-F1 on the 3-class Task 1 (9th/41) and 59 on the 9-class Task 2 (3rd/33). Across 8 transformer encoders optimized through a four-stage pipeline, partial encoder layer unfreezing outperforms full fine-tuning by a wide margin. Combining English and multilingual encoders further improves ensemble performance over either family alone, despite multilingual models being individually weaker. Prompt-based LLMs, without any task-specific parameter updates, outperform fine-tuned encoders, particularly on minority classes; among open-weight LLMs, parameter count does not predict performance. Enriched input, concatenating the full interviewer turn, improves LLM performance but not that of encoders, an effect that persists with Longformer’s extended context window, suggesting the divergence is not attributable to sequence-length capacity alone in our settings. The *Clear Reply/Ambivalent* boundary remains the dominant failure mode, mirroring the disagreement among human annotators. Our code, prompts, model configurations, and results are publicly available.¹

1 Introduction

Politicians typically avoid direct answers in interviews (Bull, 2003), hence tools for detecting evasion are valuable for political discourse research and accountability. The SemEval-2026 CLARITY task (Thomas et al., 2026) addresses the automatic detection of response clarity and evasion in English question-answer pairs from U.S. presidential interviews, building on the dataset and taxonomy

introduced by Thomas et al. (2024). Task 1 requires classifying responses into one of 3 Clarity categories, while Task 2 predicts one of 9 fine-grained Evasion labels. We participated in both tasks (with a focus on Task 1) exploring 3 approaches: (1) encoder-based models optimized through a four-stage pipeline, (2) a Longformer-based architecture to account for long-context instances, and (3) prompt-based LLMs evaluated across multiple prompting strategies with top configurations combined into a final ensemble. Due to a lack of time, the best-performing Task 1 configuration was applied to Task 2. Overall, our LLM ensemble achieves 80 macro-F1 on the Task 1 test set and 59 on the Task 2 test set, outperforming fine-tuned encoders.

2 Background

2.1 Dataset Description

The dataset provided by the organizers consists of 3,993 English question-answer pairs extracted from televised U.S. presidential interviews, split into training (3,448 instances), dev (308), and test (237) sets. Each instance is annotated at two levels: 3 Clarity labels (Level 1) and Evasion labels (Level 2) that subcategorize the Clarity labels into 9 more fine-grained categories. Figure 1 shows label distributions across both levels for the train and dev sets; train Evasion labels are color-coded by their parent Clarity class, while dev Evasion labels are shown per annotator. The training set provides a single gold label per level; while the dev set includes one gold Clarity label but 3 independent Evasion annotations per instance. The visible annotator disagreement highlights the inherent difficulty of fine-grained Evasion classification. Both levels exhibit substantial class imbalance: *Ambivalent* dominates Level 1, while Level 2 is heavily skewed toward *Explicit* and *Dodging*.

¹<https://github.com/CLaC-Lab/SemEval-2026-task6-CLARITY>




Figure 1: Label distributions across Clarity (top) and Evasion (bottom) levels for the train and dev sets.

2.2 Related Work

Political evasion has been widely studied in political science (Rasiah, 2010; Bull, 2003), with Bull (2003) finding that politicians frequently avoid directly answering interview questions. Thomas et al. (2026, 2024) formalize this phenomenon computationally through the CLARITY task², a two-level hierarchical taxonomy, and a dataset of political question-answer pairs annotated through a combination of human expertise and LLM-assisted validation, alongside baseline evaluations across multiple architectures. Building on this foundation, we investigate how encoder adaptation techniques and prompting strategies influence performance across Clarity classes, particularly under class imbalance and the hierarchical structure of the labels.

3 System Overview

We address Task 1 using 3 approaches: (1) encoder-based models, (2) Longformer-based architecture, and (3) prompt-based LLM classifiers. Among all Task 1 experiments, the LLM prompting approach with 27 shots achieved the strongest performance. Due to a lack of time, we therefore adopt this strategy for Task 2.

3.1 Encoder-Based Models

3.1.1 Models

We evaluate 8 transformer encoders spanning English and multilingual variants in base and large configurations. English models include

²<https://konstantinosftw.github.io/CLARITY-SemEval-2026/>




Figure 2: Encoder pipeline. Each model is optimized independently across stages before final ensembling.

bert-base-uncased (Devlin et al., 2019), RoBERTa-base, RoBERTa-large (Liu et al., 2019), and DeBERTa-v3-base (He et al., 2023). Multilingual counterparts include bert-base-multilingual-cased, xlm-RoBERTa-base, xlm-RoBERTa-large (Conneau et al., 2020), and mDeBERTa-v3-base. We include multilingual models to assess whether cross-lingual pretraining improves model performance.

3.1.2 Training Strategy

Figure 2 illustrates the training strategy of the encoder models where each model is optimized through 4 sequential stages. At each stage, we evaluate multiple configurations and carry forward the variant achieving the highest macro-F1. The input to all encoder models is the concatenated QA pair. To mitigate label imbalance, we apply class-weighted cross-entropy (WCE) (see Appendix A.1). We use a dropout rate (0.1) for regularization. In later stages, we experiment with alternative loss functions, dropout values and regularization strategies.

Stage 1: Encoder Adaptation. We compare (a) full fine-tuning, (b) LoRA-based tuning with a frozen encoder backbone, and (c) partial unfreezing of the top 25%, 50%, or 75% layers.

Stage 2: Classification & Representation. We evaluate 3 classification heads (default: a single lin-

ear layer over CLS; MLP over CLS; and mean pooling with a single linear layer; see Appendix B.1), each under standard truncation (preserving the first 512 tokens) and head-tail truncation (preserving the first and last 256 tokens).

Stage 3: Loss Function & Regularization. We additionally evaluate focal loss and dropout 0.3 against the WCE + dropout 0.1 baseline.

Stage 4: Data Augmentation & Input Enrichment. We evaluate two dimensions: (a) training data, comparing the original dataset with an augmented version that upsamples *Clear Non-Reply* via LLM-generated paraphrases (see Appendix D); and (b) input mode, comparing atomic input (target question only) against enriched input incorporating the full interviewer turn (see Appendix B.2).³

Ensemble Strategy. After Stage 4, we retain the configuration achieving the highest dev macro-F1 for each of the 8 encoder models and construct an ensemble using majority voting. In case of a tie, we select the label with the highest mean predicted probability score among the tied candidates.

3.2 Long-Context Model

Since approximately 30% of the training QA pairs exceed the 512-token limit, we experiment with longformer-base-4096 to assess whether long-context modelling improves performance.

Stage 1: Configuration Selection. We base our initial design on DeBERTa-v3-base, the best-performing model across all encoder stages, and we adopt its best configuration. We then conduct early exploration with alternative configurations (different classification heads, dropout 0.3, varying max lengths). Our exploration did not improve macro-F1 on the dev set; we therefore retained the original configuration.

Stage 2: Ablation Study. With the core configuration fixed, we systematically evaluate 3 design dimensions across 12 configurations: classification head type (default vs. MLP), input ordering (3 scenarios: context-question-answer, question-context-answer, question-context-answer-question-repeat), and global attention pattern (CLS-only or extended to question tokens). The maximum sequence length was set to 2048, with a 512-token attention window (see Appendix B.4).

³The full interviewer turn may contain multiple questions; we refer to this as *context* throughout.

3.3 LLM-Based Models

3.3.1 Models

We evaluate both open-weight and proprietary large language models. Open-weight families include LLaMA (LLaMA-3.1-Nemotron-Ultra-253B (Bercovich et al., 2025); LLaMA-3.3-70B-Instruct (Grattafiori et al., 2024)), Qwen (Qwen3-235B-Instruct; Qwen3-80B-Instruct; Qwen3-30B-Instruct) (Yang et al., 2025), and Mixtral (Mixtral-8x22B-Instruct). Among proprietary models, we evaluate GPT-5 (Singh et al., 2025) (OpenAI) across the full prompt sweep. After identifying the best prompting strategy, we evaluate Claude-Opus-4.5 (Anthropic) and Gemini-3-Flash-Preview (Google) under this strategy for prompt optimization and final comparison. The corresponding checkpoints for all open-weight models are provided in Appendix C.1.

3.3.2 Task 1 Prompting Strategy

Stage 1: Technique & Model Sweep. Our prompting design is informed by the survey of prompt engineering techniques of Schulhoff et al. (2025). We evaluate zero-shot (ZS), zero-shot with instruction repeated (ZS+Re2) (Xu et al., 2024), few-shot prompting with 3, 9 and 27 shots (FS3/FS9/FS27), and a prompt inspired by chain-of-thought (CoT) prompting (Kojima et al., 2022); however, we instruct the model to reason internally before producing the final label. Few-shot demonstrations are class-balanced across the 3 Level-1 labels (multiples of 3). All strategies are evaluated under both atomic and enriched input settings. Prompt templates are provided in Appendix C.2.

Stage 2: Prompt Optimization & Model Expansion. Building on FS27 with enriched input (target question concatenated with the full interviewer turn), the best-performing prompt during Stage 1, we conduct prompt optimization by introducing two controlled modifications: (a) prepending the president’s name, and (b) augmenting each Clarity label definition with its Level-2 Evasion subcategories, while still predicting only the Level-1 label. The refined prompt is then extended to Claude-Opus-4.5 and Gemini-3-Flash-Preview to assess cross-model robustness.

Ensemble Strategy. After prompt optimization and model expansion, we select the top 3 configurations based on dev macro-F1: GPT-5, Gemini-3-Flash-Preview, and Qwen3-235B, all

under the FS27 enriched setting with subcategory-augmented definitions. Gemini and Qwen additionally incorporate president-name conditioning, while GPT-5 does not as it performed better without it. Predictions are combined via majority voting, defaulting to *Ambivalent* on ties.

3.3.3 Task 2 Prompting Strategy

For Task 2, we did not perform a separate prompt sweep due to a lack of time. Instead, we used the best-performing ensemble configuration from Task 1: GPT-5, Gemini-3-Flash-Preview, and Qwen3-235B-Instruct, all under FS27 using enriched input and subcategory-augmented definitions. Gemini and Qwen additionally incorporate president-name conditioning, while GPT does not. Predictions are combined via majority voting. In case of ties, we default to the class with the highest frequency among the tied classes in the training set.

4 Experimental Setup

Encoder and Longformer models are implemented in PyTorch using Hugging Face Transformers. All models are trained on the official training split and selected based on dev macro-F1. Class imbalance is addressed via weighted cross-entropy (Appendix A.1). LLM-based experiments are conducted via official APIs, with prompt templates and few-shot demonstrations held fixed across models. Complete hyperparameters, hardware setup, Longformer configurations, and prompt templates are provided in Appendices A.2, B.4, and C.2.

5 Results & Analysis

Table 2 reports Task 1 dev macro-F1 for each encoder using its best configuration after the four-stage optimization process (see Section 3.1.2). For LLM-based systems, we report the top-performing models retained for ensembling after the two-stage prompting process (see Appendix C.3 for the best configuration of the remaining LLMs). Among encoder-based models, DeBERTa-v3-base performs best (65.1), and ensembling the 8 optimized encoders improves performance to 70.5 on the dev set; while the best-performing Longformer configuration achieves 64.3. LLM-based systems outperform fine-tuned encoders, with Gemini-3-Flash-Preview reaching 71.9. The LLM ensemble, used for both our Task 1 and Task 2 submissions, achieves 78.1 on the dev set and 80 on the

Design Choice	Avg F1 (Δ)	# Sel.
Stage 1: Encoder Adaptation		
<i>Full Finetuning</i>	41.2	0
Unfreeze Top 25%	59.1 (+17.9)	6
LoRA ($r = 16$)	56.1 (+14.9)	1
Unfreeze Top 75%	45.1 (+3.9)	1
Stage 2: Classification & Representation		
<i>Def. Head + Std. Trunc.</i>	60.1	2
Def. Head + Head-Tail Trunc.	59.9 (-0.2)	2
MeanPool + Head-Tail Trunc.	59.9 (-0.2)	1
MeanPool + Std. Trunc.	59.8 (-0.3)	2
MLP Head + Head-Tail Trunc.	57.8 (-2.3)	1
Stage 3: Loss & Regularization		
<i>WCE + Dropout = 0.1</i>	61.5	8
Stage 4: Data Augmentation & Input Enrichment		
<i>Original Data + Atomic Input</i>	61.5	5
Augmented Data + Enriched Input	59.8 (-1.7)	2
Augmented Data + Atomic Input	59.8 (-1.7)	1
Original Data + Enriched Input	59.1 (-2.4)	0

Table 1: Average dev macro-F1 across 8 encoder models per design choice. Δ is relative to the stage baseline, and # Sel. is the number of best models selecting that choice. *Italicized rows* indicate the stage baseline.

System Type	Model	Macro-F1
Encoder-Based	DeBERTa-v3-base	65.1
	RoBERTa-base	63.8
	BERT-base	63.4
	RoBERTa-large	63.0
	mBERT-base	62.9
	xlm-RoBERTa-base	62.6
	xlm-RoBERTa-large	61.4
	mDeBERTa-v3-base	60.8
	Ensemble (EN only)	65.6
	Ensemble (Multi only)	68.2
	Ensemble (EN+Multi)	70.5
Long Context	Longformer-base-4096	64.3
LLM-Based	Gemini-3-Flash-Preview	71.9
	GPT-5	71.5
	Qwen3-235B-Instruct	68.3
	LLM Ensemble	78.1

Table 2: Task 1 dev macro-F1 across system approaches. Each row reports the best config after full optimization.

test set. Applying this configuration to Task 2 yields 59 macro-F1 on the test set.

5.1 Encoder and Long-Context Results

Table 1 reports, for each encoder design choice, the average dev macro-F1 across the 8 encoder models, its change (Δ) relative to the stage baseline, and the number of models that selected that choice in their optimal configuration. For each stage, the baseline is the average dev macro-F1 of the best configuration carried over from all previous stages, except for Stage 1, where full fine-tuning serves as the baseline.

Table 1 shows that encoder adaptation is a critical optimization step: partially unfreezing the top 25% layers improves macro-F1 by +17.9 over full




Figure 3: Impact of prompting strategy, input configuration, and LLM choice on dev macro-F1. Each panel reports averages over the remaining variables, computed across all experimental configurations.

fine-tuning (59.1 vs. 41.2), suggesting that full fine-tuning may overfit the limited training size. Subsequent stages build on the best configuration of the previous stage but yield only marginal changes. In Stage 2, no classification head or truncation strategy consistently improves across all models over the carried-forward baseline; in Stage 3, all 8 models selected weighted cross-entropy with dropout 0.1 as their optimal configuration, indicating that class weighting is sufficient without additional loss reshaping. In Stage 4, data augmentation decreases average performance across models (61.5 to 59.8 macro-F1). However, its effect is model-dependent: 3 models, DeBERTa-v3-base, mBERT-base, and xlm-RoBERTa-base, improve under at least one augmented configuration (see Appendix B.3), suggesting that paraphrase-based augmentation does not generalize uniformly across architectures.

As shown in Table 2, ensembling the 8 optimized encoders yields an additional +5.4 macro-F1 (65.1 to 70.5), indicating that architectural diversity provides complementary signal. While English-only encoders achieve higher standalone macro-F1 than multilingual variants, the multilingual-only ensemble (68.2) surpasses the English-only ensemble (65.6), and combining both families produces the strongest result (70.5). This suggests that multilingual pretraining produces models with distinct decision boundaries, which improve overall performance when combined with English-only encoders. The Longformer (64.3) fell below the encoder ensemble despite its extended context capacity (see Appendix B.4 for further details).

Ensemble	Clear Reply	Ambivalent	Clear Non-Reply
LLM	70.7	85.3	78.4
Encoder	61.0	82.4	68.1

Table 3: Dev F1 per class for the encoder and LLM ensembles.

5.2 LLM-Based Results

Figure 3 summarizes dev macro-F1 results from the Stage 1 prompting sweep, prior to prompt refinement and model expansion. Each panel reports averages over the remaining variables, computed across all experimental configurations. The top row isolates design choices, comparing prompting strategy (left) and input configuration (right). Few-shot prompting (FS27) achieves the strongest average performance (56.4), while the enriched input consistently outperforms atomic input (+4.7), indicating consistent gains across configurations from inference-time design choices. The bottom panel shifts focus to LLM choice, revealing larger variation across models: GPT-5 reaches 62.0 macro-F1, whereas LLaMA-3.1-Nemotron-Ultra-253B trails at 31.5. Notably, parameter count alone does not explain these differences. While having additional context leads to measurable improvements, the choice of LLM accounts for substantially larger performance variation in Clarity detection. The best configuration for each model is reported in Appendix C.3.

Prompt refinement in Stage 2 further improves performance across models. For GPT-5, macro-F1 rises from 64.4 under the best Stage 1 configuration to 71.5 after refinement (+7.1), driven by augmenting the 3 Clarity label definitions with their corresponding Level 2 Evasion subcategory descriptions. This result highlights the impact of prompt engineering and hierarchical label information on performance, without any model fine-tuning.

5.2.1 Cross-System Comparison

Table 3 shows dev F1 per class for the encoder and LLM ensembles. The LLM ensemble improves performance across all classes, with the largest gains on the minority classes, *Clear Reply* (+9.7) and *Clear Non-Reply* (+10.3). This pattern suggests greater robustness to class imbalance in underrepresented categories. The enriched input reveals system divergence; for the encoder models, adding the full interviewer turn reduced the average macro-F1 (e.g., Original + Enriched: 59.1, $\Delta = -2.4$; see Table 1). The Longformer, de-

spite long-context modeling, underperforms standard encoders in our settings. In contrast, LLMs consistently gain from additional context (54.1 vs. 49.4, $\Delta = +4.7$; Figure 3), suggesting that LLMs leverage enriched context more effectively than encoder architectures, beyond sequence-length capacity alone. Error analysis (Appendix E) reveals that both ensembles share the same dominant error pattern: bidirectional confusion between *Clear Reply* and *Ambivalent*, mirroring the lowest pairwise annotator agreement among Clarity classes ($\kappa = 0.65$; Thomas et al. 2024). The LLM ensemble improves *Clear Non-Reply* recall (20/23 vs. 16/23), consistent with its stronger minority class performance seen in Table 3.

6 Conclusion

We systematically compared encoder adaptation and prompt-based LLMs for SemEval-2026 CLARITY. Our LLM ensemble achieves 80 macro-F1 on Task 1 (9th/41) and 59 on Task 2 (3rd/33), with prompt refinement and hierarchical label information driving the largest gain. Partial unfreezing improves encoder performance, yet inference-time prompting proves more effective than fine-tuning, particularly on minority classes. Ensemble diversity matters more than model scale, with multilingual models contributing better when combined with English-only encoders. The persistent *Clear Reply/Ambivalent* confusion mirrors annotator disagreement, reflecting the inherent ambiguity of political discourse. Together, these findings suggest that for low-resource, hierarchically-labeled tasks like CLARITY, inference-time design choices (prompt structure, label hierarchy, ensemble composition) are more impactful than parameter-level adaptation.

Limitations and Future Work

The dataset is restricted to U.S. presidential interviews and contains fewer than 4K instances, which may limit generalization to other political settings or languages. The boundary between *Clear Reply* and *Ambivalent* is inherently subjective, which is reflected in moderate annotator agreement ($\kappa = 0.65$; Thomas et al. 2024). LLM-based results rely partially on proprietary APIs (GPT, Gemini), which may affect long-term reproducibility due to model versioning. Our CoT-inspired prompting prevents inspection of intermediate reasoning, limiting interpretability and may have affected results. Our

encoder study spans 136 configurations across 8 models and 4 optimization stages; due to computational constraints, each configuration was evaluated once with a fixed random seed, and variance across seeds was not assessed. Future work should quantify this variance through multi-seed runs on the final per-model configurations to provide stability estimates for the reported encoder results.

Several findings in this paper are reported empirically but not yet causally explained: *why* multilingual encoders underperform individually but strengthen the ensemble, *why* enriched input helps LLMs but degrades encoder performance, and *why* Longformer fails to leverage its long-context capacity in this setting. We frame the latter as observed under our current configurations rather than a general property of long-context models.

Our Task 2 system reuses the optimal Task 1 ensemble configuration without an independent prompt sweep over the 9-class label space. While this transfer resulted in competitive performance (3rd/33), it does not establish that the configuration is optimal for the Evasion classification task. Future work should conduct a dedicated Task 2 prompt sweep to determine whether the Task 1 optimum transfers to the 9-class Evasion classification, or whether fine-grained Evasion classification benefits from different prompting strategies.

Finally, we did not perform a head-to-head comparison with the top-ranked Task 6 systems (TeleAI at 89 macro-F1 on Task 1). Understanding whether the 9-point gap stems from prompt design, model choice, or fundamentally different architectures would contextualize our results; we leave such cross-system analysis to future work once participating systems are publicly described.

Acknowledgments

We thank the SemEval-2026 Task 6 organizers for designing and running the CLARITY shared task, and the anonymous reviewers for their constructive feedback. We also thank the Pierre Arbour Foundation, the Natural Sciences and Engineering Research Council of Canada (NSERC), and the Fonds de recherche du Québec (FRQ) for their financial support.

References

- Akhiad Bercovich, Itay Levy, Izik Golan, Mohammad Dabbah, and 132 others. 2025. [Llama-Nemotron: Efficient Reasoning Models](#). *Preprint*, arXiv:2505.00949.
- Peter Bull. 2003. *The Microanalysis of Political Communication: Claptrap and Ambiguity*. Routledge.
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. [Unsupervised Cross-lingual Representation Learning at Scale](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8440–8451, Online. Association for Computational Linguistics.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, and 557 others. 2024. [The Llama 3 Herd of Models](#). *Preprint*, arXiv:2407.21783.
- Pengcheng He, Jianfeng Gao, and Weizhu Chen. 2023. [DeBERTaV3: Improving DeBERTa using ELECTRA-Style Pre-Training with Gradient-Disentangled Embedding Sharing](#). In *Proceedings of Eleventh International Conference on Learning Representations*.
- Takeshi Kojima, Shixiang (Shane) Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. [Large Language Models are Zero-Shot Reasoners](#). In *Advances in Neural Information Processing Systems*, volume 35, pages 22199–22213. Curran Associates, Inc.
- Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. [RoBERTa: A Robustly Optimized BERT Pretraining Approach](#). *Preprint*, arXiv:1907.11692.
- Parameswary Rasiiah. 2010. [A Framework for the Systematic Analysis of Evasion in Parliamentary Discourse](#). *Journal of Pragmatics*, 42(3):664–680.
- Sander Schulhoff, Michael Ilie, Nishant Balepur, Konstantine Kahadze, and 27 others. 2025. [The Prompt Report: A Systematic Survey of Prompt Engineering Techniques](#). *Preprint*, arXiv:2406.06608.
- Aaditya Singh, Adam Fry, Adam Perelman, Adam Tart, and 480 others. 2025. [OpenAI GPT-5 System Card](#). *Preprint*, arXiv:2601.03267.
- Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2024. [“I Never Said That”: A Dataset, Taxonomy and Baselines on Response Clarity Classification](#). In *Findings of the Association for Computational Linguistics (EMNLP 2024)*, pages 5204–5233, Miami, Florida, USA. Association for Computational Linguistics.
- Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2026. [SemEval-2026 Task 6: CLARITY – Unmasking Political Question Evasions](#). *Preprint*, arXiv:2603.14027.
- Xiaohan Xu, Chongyang Tao, Tao Shen, Can Xu, Hongbo Xu, Guodong Long, Jian-Guang Lou, and Shuai Ma. 2024. [Re-Reading Improves Reasoning in Large Language Models](#). In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP 2024)*, pages 15549–15575, Miami, Florida, USA. Association for Computational Linguistics.
- An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, and 56 others. 2025. [Qwen3 Technical Report](#). *Preprint*, arXiv:2505.09388.

A Optimization and Loss

A.1 Loss Functions

To address label imbalance, we assign each class i a weight w_i computed as:

$$w_i = \frac{N}{C \cdot n_i}$$

where N is the total number of training instances; C , the number of classes; and n_i , the number of instances in class i . This places greater weight on rarer classes during training. All weights are computed from the training split only. Both loss functions evaluated in Stage 3 incorporate these class weights.

Weighted Cross-Entropy (WCE). The standard formulation used as the default loss across all stages, applying class weights to the cross-entropy term:

$$\mathcal{L}_{\text{WCE}} = -w_t \log(p_t)$$

where p_t denotes the predicted probability of the ground-truth class t .

Focal Loss. Evaluated as an alternative in Stage 3, this formulation adds a focal factor $(1 - p_t)^\gamma$ that down-weights well-classified examples to focus training on harder instances:

$$\mathcal{L}_{\text{Focal}} = -w_t(1 - p_t)^\gamma \log(p_t)$$

where $\gamma = 2.0$.

A.2 Training and Decoding Configuration

Table 4 reports the training configuration for encoder and Longformer models. For partial unfreezing, the embedding layer and all but the top $k\%$ of encoder layers are frozen; the classification head remains trainable throughout. For LLM-based experiments, Hugging Face-hosted models use temperature 0; all other models (GPT, Gemini, Claude) use default API generation settings. All encoder and Longformer experiments were conducted on an NVIDIA RTX 4070 12GB with 1 GPU. Training time per model ranged from approximately 5 minutes to 6 hours. Software dependencies are listed in Table 5.

Parameter	Encoder	Longformer
Optimizer	AdamW	AdamW
Learning rate	5×10^{-5}	2×10^{-5}
Weight decay	0.01	0.01
Batch size	16	2
Gradient accumulation	–	8
Effective batch size	16	16
Max epochs	20	20
Early stopping patience	5	7
Max sequence length	512	2048
Attention window	–	512
Dropout	0.1*	0.1
Partial unfreezing	top $k\%$	top 25%

Table 4: Training configuration for encoder and Longformer models. *Dropout 0.3 is additionally evaluated in Stage 3.

Dependency	Version
Python	3.12.7
PyTorch	2.5.1+cu118
Transformers	4.51.3
NumPy	2.3.5
pandas	2.2.3

Table 5: Software dependencies used in experiments.

B Architecture and Model Variants

B.1 Classification Heads

We evaluate 3 classification head architectures:

- **Default:** CLS $\rightarrow$ Dropout(0.1) $\rightarrow$ Linear.
- **MLP:** CLS $\rightarrow$ Dropout(0.1) $\rightarrow$ Linear $\rightarrow$ GELU $\rightarrow$ Dropout(0.1) $\rightarrow$ Linear.
- **Average Pooling:** Mean pooling over non-padding token embeddings $\rightarrow$ Linear.

B.2 Enriched Input Configuration

In the enriched input setting, we concatenate the target question with context (i.e., the full interviewer turn) using explicit textual markers:

Target question: *<question>*

Full interviewer turn (context): *<interview_question>*

This concatenated string forms the first sequence input to the encoder. The interview answer is provided as the second sequence.

B.3 Encoder Final Configurations

The final configuration per encoder model after the four-stage optimization is as follows, ordered by dev macro-F1. All models selected weighted cross-entropy (WCE) loss and a dropout rate of 0.1 as part of their optimal configuration.

1. DeBERTa-v3-base (macro F1: 65.1): Unfreeze 25%, MLP head, head-tail truncation, augmented data, enriched input.
2. RoBERTa-base (macro F1: 63.8): Unfreeze 75%, mean pooling head, head-tail truncation, original data, atomic input.
3. BERT-base (macro F1: 63.4): LoRA (rank 16), default head, head-tail truncation, original data, atomic input.
4. RoBERTa-large (macro F1: 63.0): Unfreeze 25%, mean pooling head, standard truncation, original data, atomic input.
5. mBERT-base (macro F1: 62.9): Unfreeze 25%, default head, standard truncation, augmented data, enriched input.
6. xlm-RoBERTa-base (macro F1: 62.6): Unfreeze 25%, default head, standard truncation, augmented data, atomic input.
7. xlm-RoBERTa-large (macro F1: 61.4): Unfreeze 25%, default head, head-tail truncation, original data, atomic input.
8. mDeBERTa-v3-base (macro F1: 60.8): Unfreeze 25%, mean pooling head, standard truncation, original data, atomic input.

B.4 Longformer Configuration & Ablation

We evaluate 12 configurations across 3 design dimensions: input ordering, global attention pattern, and classification head type. All configurations use max sequence length 2048 and attention window 512. Figure 4 reports mean dev macro-F1 per design choice, averaged over the remaining dimensions. Error bars represent one standard deviation across experiments sharing the same value of the given configuration dimension. The overlapping error bars across all 3 dimensions indicate that no single design choice yields a reliable improvement, and the results should be interpreted with caution. The results do not indicate a clear superior design choice, and further experimentation is needed.

Global Atten.	Head	Input Order	Dev F1
CLS	MLP	C+Q+A	64.3
CLS+Q	MLP	C+Q+A	62.6
CLS+Q	Default	Q+C+A+Q	62.1
CLS	MLP	Q+C+A	61.9
CLS	Default	C+Q+A	61.4
CLS+Q	MLP	Q+C+A+Q	61.1
CLS+Q	MLP	Q+C+A	60.8
CLS	Default	Q+C+A	60.5
CLS	MLP	Q+C+A+Q	60.4
CLS+Q	Default	Q+C+A	60.2
CLS	Default	Q+C+A+Q	60.1
CLS+Q	Default	C+Q+A	57.2

Table 6: Longformer ablation across 12 configurations, ordered by dev macro-F1. Input order reflects the token sequence fed to the model: C = full interviewer turn (context), Q = target question, A = answer. CLS+Q extends global attention to question tokens in addition to the CLS token.




Figure 4: Mean dev macro-F1 per design choice for the Longformer ablation. Error bars represent one standard deviation across configurations sharing each value.

C LLM Details

C.1 Model Checkpoints

We report the exact Hugging Face checkpoints used for all LLM experiments to ensure reproducibility (Table 7).

Model	Checkpoint
LLaMA-3.1-Nemotron Ultra-253B	Llama_3_1_Nemotron_Ultra_253B_v1
LLaMA-3.3-70B-Instruct	Llama-3.3-70B-Instruct
Qwen3-235B-Instruct	Qwen3-235B-A22B-Instruct-2507
Qwen3-80B-Instruct	Qwen3-80B-A3B-Instruct
Qwen3-30B-Instruct	Qwen3-30B-A3B-Instruct-2507
Mixtral-8x22B-Instruct	Mixtral-8x22B-Instruct

Table 7: Open-weight large language models (LLMs) used in our experiments, with their corresponding checkpoints.

C.2 Prompt Templates

All Stage 1 prompting strategies (see Section 3.3.2) share a common system role and label definitions, shown in the Zero-Shot (ZS) template below. Subsequent strategies extend this base prompt (ZS+Re2 repeats the instruction block; FS prepends demonstrations; CoT-inspired prompts add internal reasoning instructions). Stage 2 prompt optimization refines the best-performing Stage 1 configuration (FS27) by expanding Clarity label definitions with Level-2 (i.e., Evasion) subcategory descriptions and, in some variants, prepending the president’s name. Input placeholders {Q} and {A} are replaced at inference time with the target question and answer, respectively. In the enriched input setting, {Q} is replaced with the concatenated target question and full interviewer turn as described in Appendix B.2.

Zero-Shot (ZS). Defines the role, label set, and output format. All subsequent strategies extend this base prompt.

You are a world-class political discourse analyst trained to detect evasive communication strategies in high-stakes interviews.

Judge the clarity of an answer to a journalist’s question.

Labels (choose exactly ONE):

- Clear Reply: the answer directly supplies what was asked.
- Ambivalent Reply: the answer references the topic but is indirect, vague, partial, or hedged.

- Clear Non-Reply: the answer refuses, claims not to know, asks for clarification, or ignores the question.

Output ONLY the label name (no explanation, no punctuation).

QUESTION: {Q}
ANSWER: {A}
Label:

Zero-Shot with Instruction Repetition (ZS+Re2).

Extends ZS by repeating the instruction block before the input.

You are a world-class political discourse analyst trained to detect evasive communication strategies in high-stakes interviews.

Judge the clarity of an answer to a journalist's question.

Labels (choose exactly ONE):

- Clear Reply: the answer directly supplies what was asked.
- Ambivalent Reply: the answer references the topic but is indirect, vague, partial, or hedged.
- Clear Non-Reply: the answer refuses, claims not to know, asks for clarification, or ignores the question.

Output ONLY the label name (no explanation, no punctuation).

Read the question again:

Judge the clarity of an answer to a journalist's question.

Labels (choose exactly ONE):

- Clear Reply: the answer directly supplies what was asked.
- Ambivalent Reply: the answer references the topic but is indirect, vague, partial, or hedged.
- Clear Non-Reply: the answer refuses, claims not to know, asks for clarification, or ignores the question.

Output ONLY the label name (no explanation, no punctuation).

QUESTION: {Q}
ANSWER: {A}
Label:

Few-Shot (FS). Few-shot prompts extend the ZS base by inserting k labelled demonstrations between the label definitions and the input. Demonstrations are class-balanced (equal number per class) and drawn from the training set. We evaluate $k \in \{3, 9, 27\}$ shots. The structure is as follows:

You are a world-class political discourse analyst trained to detect evasive communication strategies in high-stakes interviews.

Judge the clarity of an answer to a journalist's question.

Labels (choose exactly ONE):

- Clear Reply: the answer directly supplies what was asked.
- Ambivalent Reply: the answer references the topic but is indirect, vague, partial, or hedged.
- Clear Non-Reply: the answer refuses, claims not to know, asks for clarification, or ignores the question.

The following examples illustrate each label:

Example 1:

Question: {example_question}
Answer: {example_answer}
Label: {label}

[k class-balanced demonstrations]

Output ONLY the label name (no explanation, no punctuation).

Question: {Q}
Answer: {A}
Label:

CoT-Inspired (CoT). The CoT-inspired prompt replaces the output instruction of the ZS base with an internal reasoning directive. Unlike standard CoT, the model is instructed to reason internally without outputting any intermediate steps:

You are a world-class political discourse analyst trained to detect evasive communication strategies in high-stakes interviews.

Judge the clarity of an answer to a journalist's question.

Labels (choose exactly ONE):

- Clear Reply: the answer directly supplies what was asked.
- Ambivalent Reply: the answer references the topic but is indirect, vague, partial, or hedged.
- Clear Non-Reply: the answer refuses, claims not to know, asks for clarification, or ignores the question.

In your own mind, reason step by step about how the answer responds to the question. Keep all intermediate reasoning hidden and do not write it out.

When you have decided, output ONLY the label name (no explanation, no extra text, no punctuation).

QUESTION: {Q}

ANSWER: {A}

Label:

Stage 2: Subcategory-Augmented Optimized Prompt. Building on FS27, Stage 2 replaces the Clarity label definitions with expanded versions that include Level-2 Evasion subcategories. The user message includes the target question, the context (full interviewer turn), and the answer, and optionally prepends the speaker's name.

You are a world-class political discourse analyst trained to detect evasive communication strategies in high-stakes interviews.

Judge the clarity of an answer to a journalist's question.

Labels (choose exactly ONE):

Ambivalent Reply

Definition: Where a response is given in the form of a valid answer but allows for multiple interpretations.

Sub-categories:

- Implicit: The information requested is given, but without being explicitly stated (not in the expected form)
- General: The information provided is too general/lacks the requested specificity
- Partial: Offers only a specific component of the requested information
- Dodging: Ignoring the question altogether
- Deflection: Starts on topic but shifts the focus and makes a different point than what is asked

Clear Reply

Definition: Containing replies that admit only one interpretation.

Sub-category:

- Explicit: The information requested is explicitly stated (in the requested form)

Clear Non-Reply

Definition: Containing responses where the answerer openly refuses to share information.

Sub-categories:

- Declining to answer: Acknowledge the question but directly or indirectly refusing to answer at the moment
- Claims ignorance: The answerer claims/admits not to know the answer themselves
- Clarification: Does not provide the requested information and asks for clarification

The following examples illustrate each category:

[27 class-balanced demonstrations]

Output ONLY the category name.

Target question (to evaluate): {Q}
Speaker: {president}
Full interviewer turn (may contain multiple questions): {context}
Answer: {A}
Label:

Task 2: Few-Shot Evasion Classification (FS27)

Task 2 mirrors the optimized Task 1 Stage 2 prompt (FS27, enriched input, subcategory-augmented definitions), adapted to the 9 Evasion categories. Demonstrations are class-balanced across the 9 Evasion categories and drawn from the training split. The structure is as follows:

You are a world-class political discourse analyst trained to detect evasive communication strategies in high-stakes interviews.

Your task is to classify how a public official answers a journalist's question.

Labels (choose exactly ONE):

- Implicit: The information requested is given, but without being explicitly stated (not in the expected form).
- General: The information provided is too general or lacks the requested specificity.
- Partial/half-answer: Offers only a specific component of the requested information.
- Dodging: Ignoring the question altogether.
- Deflection: Starts on topic but shifts the focus and makes a different point than what is asked.
- Explicit: The information requested is explicitly stated (in the requested form).

- Declining to answer: Acknowledges the question but directly or indirectly refuses to answer.
- Claims ignorance: The answerer claims or admits not to know the answer.
- Clarification: Does not provide the requested information and asks for clarification.

The following examples illustrate each category:

[27 class-balanced demonstrations]

Output ONLY the category name.

Target question (to evaluate): {Q}
Speaker: {president}
Full interviewer turn (may contain multiple questions): {context}
Answer: {A}
Label:

C.3 Stage 1: Technique & Model Sweep: Best Configurations per LLM

Table 8 reports the best Stage 1 configuration per LLM on the dev set, ordered by dev macro-F1.

Model	Strategy	Dev F1
Qwen3-235B-Instruct	FS27	67.2
Qwen3-80B-Instruct	FS3	64.8
GPT-5	FS27	64.4
LLaMA-3.3-70B-Instruct	FS27	62.5
Mixtral-8x22B-Instruct	ZS+Re2	58.5
Qwen3-30B-Instruct	FS27	58.1
LLaMA-3.1-Nemotron-Ultra-253B	FS27	43.5

Table 8: Best Stage 1 configuration per LLM. All models use enriched input.

C.4 Stage 2: Prompt Optimization & Model Expansion: Model Configurations

The two top Stage 1 models were both Qwen variants (see Table 8, Stage 1 results). To ensure model family diversity, we retain Qwen3-235B (67.2) and GPT-5 (64.4) and expand to Gemini-3-Flash-Preview and Claude-Opus-4.5. All models use FS27 with enriched input and subcategory-augmented label definitions. The top 3 after Stage 2 (Table 9) were chosen to form the LLM ensemble; Claude-Opus-4.5 was not among the top 3 and was not retained.

Model	President	Dev F1
Gemini-3-Flash-Preview	Yes	71.9
GPT-5	No	71.5
Qwen3-235B-Instruct	Yes	68.3

Table 9: Top 3 Stage 2 results per LLM (FS27, enriched input, subcategory-augmented label definitions). **President** indicates whether the president’s name was included in the prompt. These 3 models form the final ensemble.

D Data Augmentation

D.1 Augmentation Procedure

To mitigate class imbalance, as part of Stage 4, we augment the minority class in Task 1 (*Clear Non-Reply*) by generating paraphrases using GPT-4o (temperature = 0.55). Starting from 356 *Clear Non-Reply* instances in the training set, we generate one paraphrase per instance (356 total), followed by an additional 194 paraphrases sampled uniformly at random from the same instances, yielding 550 new candidates in total. This oversampling accounts for expected filtering losses. After post-generation filtering (see Appendix D.3), 356 valid paraphrases (from the 550 new instances) are retained and appended to the training set, increasing the minority class from 356 to 712 instances.

D.2 Prompt-Level Constraints

The paraphrase generation prompt imposes a set of structural and behavioral constraints. Each rewritten answer must remain a *Clear Non-Reply*: it must not directly address the interviewer’s question and must retain its defective, no-answering intent. The prompt further requires preserving the original tone, level of formality, and stylistic markers (e.g., vagueness, hesitation); avoiding the introduction of new facts, examples, or topics; and maintaining the original discourse structure without summarizing, clarifying, or restructuring the response.

D.3 Post-Generation Filtering

Each generated paraphrase is kept only if it satisfies both conditions relative to its source: (1) the total word count within $\pm 20\%$ of the original, and (2) the sentence count within ± 1 of the original. Generation of paraphrases is retried up to 3 times per instance. Only paraphrases passing both filters are eligible for inclusion in the final training set. From the filtered pool, the first 356 valid paraphrases are appended to the original training data, doubling the size of the minority class.

E Confusion Matrices

Both the encoder and LLM ensembles (Figures 5 and 6) share the same dominant error pattern: bidirectional confusion between *Clear Reply* and *Ambivalent* (28 + 31 = 59 instances for the encoder, 21 + 27 = 48 for the LLM ensemble). This mirrors the lowest pairwise annotator agreement among Clarity classes ($\kappa = 0.65$; Thomas et al. 2024), suggesting this boundary is inherently difficult for both humans and models. The LLM ensemble improves *Clear Non-Reply* recall (20/23 vs. 16/23), consistent with its stronger minority class performance reported in Table 3.




Figure 5: Encoder ensemble confusion matrix (dev set).




Figure 6: LLM ensemble confusion matrix (dev set).